

Martin VASSOR

DSLab, EPFL



January 8, 2017

Outline

Introduction

Implementation

- Modifications

- Performance evaluation

- Performance results

Verification

- What to prove ?


- Proof steps

Conclusion


- Hash Table software

- Side effects


Naive hash table




Naive hash table




Naive hash table




Double hashing



Double hashing



Double hashing




Provided implementation

`findEmpty, findKey`

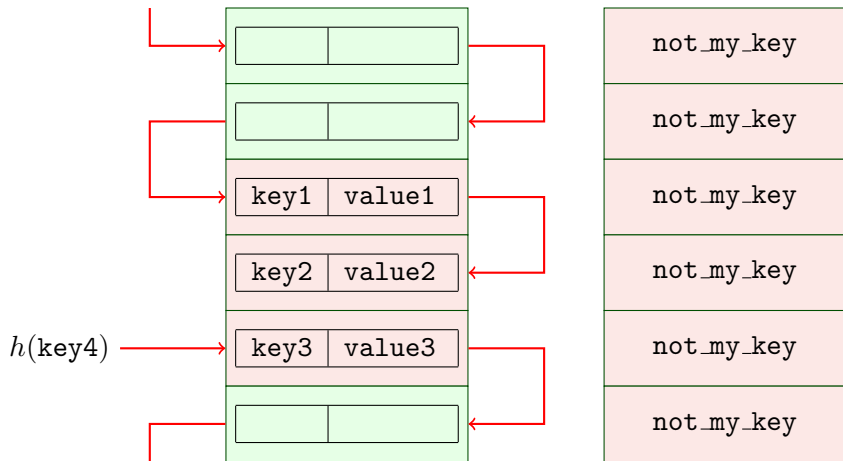
Provided verification

Example: successful search of `key3`



Provided verification

Example: unsuccessful search of key4



Provided verification

For insertion:

- ▶ Same idea
- ▶ Property: `findEmpty`

Outline

Introduction

Implementation

- Modifications

- Performance evaluation

- Performance results

Verification

- What to prove ?

- Proof steps

Conclusion

- Hash Table software

- Side effects

Modifications

- ▶ 64 bits hashes.

offset	entry
--------	-------

Except type changes, only **for** loops modified.

Performance evaluation


- ▶ Build a benchmark tool.
- ▶ Size, number of accesses, load, read/write ratio, etc...
- ▶ Converter to C file.
- ▶ First warms-up, then measures when target load is reached.

```
test_load.sh length read_ratio load1 [load2...]
```

Evaluation cases


- ▶ Worst case: searching a non existing element.
1. Allow searching non existing element.
 2. Search only existing element.

Result



Result – only existing

load results (70% read operations, 1000ops, access only existing elements)



Outline

Introduction

Implementation

- Modifications

- Performance evaluation

- Performance results

Verification

- What to prove ?

- Proof steps

Conclusion

- Hash Table software

- Side effects

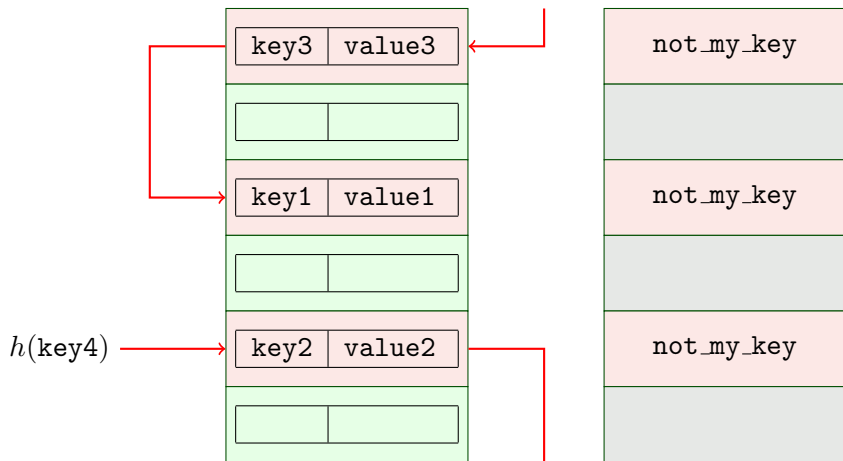
Goal

Goal: show that increment by offset covers all the map.

- ▶ Not always true (chinese remainder theorem).
- ▶ Requires: **offset** and **capacity** coprime ($\gcd = 1$).

Goal

Insert **key4**: $h_1(\text{key4}) = 5$, $h_2(\text{key4}) = 2$.



i++i

Outline

Introduction

Implementation

- Modifications

- Performance evaluation

- Performance results

Verification

- What to prove ?

- Proof steps

Conclusion

- Hash Table software

- Side effects

Hash-Table software

- ▶ Efficient (when key is present).
- ▶ Formally verified.
- ▶ Requires `capacity` and `offset` coprime.

Side effects

- ▶ 6 commits in Verifast tree (`long long` support).
- ▶ 9 issues on Verifast.
- ▶ A random access sequence generator & benchmark.

Q&A