

ÉCOLE POLYTECHNIQUE FÉDÉRALE DE LAUSANNE

OPTIONNAL SEMESTER PROJECT

Verified double-hashing hash map

Martin VASSOR

supervised by
G. CANDEA, Pr.
K. ARGYRAKI, Pr.
A. ZAOSTROVNYKH, Ph. D. Student

January 11, 2017

1 Introduction

2 Implementation

3 Verification

3.1 Provided proof

3.1.1 Requirement (R)

Figure 1 present the relevant part of the proof of the original loop in `find_key`. The last statement (`no_key_found(ks, k)`) requires the property `not_my_key(k)` to be verified for all current keys in the mapping, ensured by the `up_to(nat_of_int(length(ks)), ... (not_my_key)(k) ...)` statement.

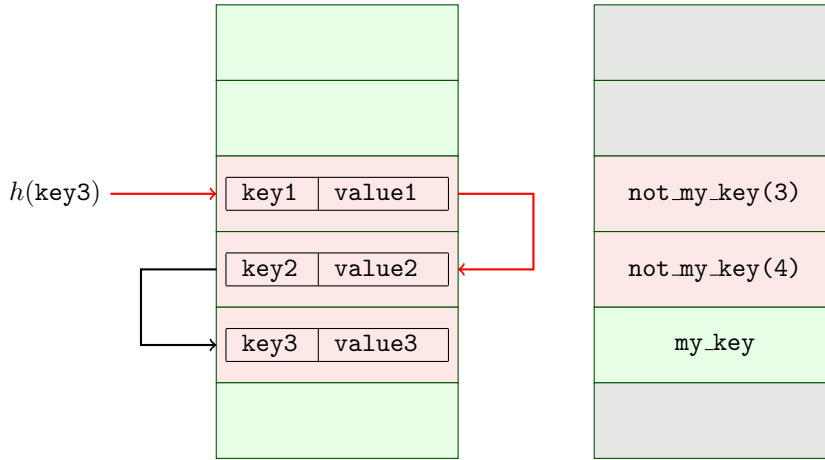
This `up_to` statement is proved by the *for*-loop invariant: at each round, `not_my_key(k, nth(index, ks))` is ensured, either because the cell is empty (`no_busy_no_key` lemma), either because the hash does not match (`no_hash_no_key` lemma), either because the key does not match (hence inferred by Verifast).

Finally, the *for*-loop only proves that the `up_to` statement holds when starting from `index = start` and looping. The lemma `by_loop_for_all` prove that this loop access is equivalent to a continuous access from 0 to `length`.

Example 1 represent a successful search in the map. The search starts at index `h(key3)`. As long as the key is not found at index `i`, `not_my_key(i)` is asserted. Finally, when `key3` is found, it is ensured to be the right key and returned.

`up_to(nat, prop)`
verifies that `prop`
is ensured for all `i`
below `nat`:
`up_to(0, prop) =`
`true`
`up_to(n, prop) =`
`prop(n-1) &&`
`up_to(n-1, prop)`

Example 1: Successful search

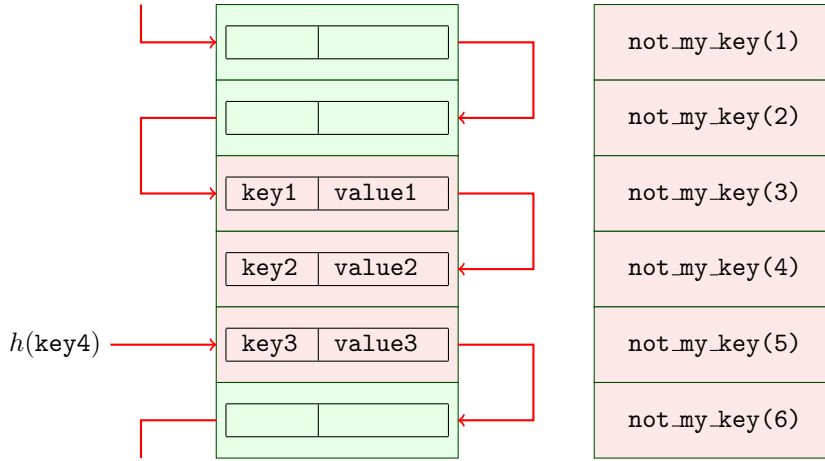


In case of unsuccessful search, as in Example 2, `not_my_key(i)` is asserted for all indexes, ensuring that the key is not present in the map.

Hence, an invariant of the *for*-loop is that `not_my_key` is asserted *for all indexes from start up to i*, by ring accesses, *i* being the loop iterator.

Example 2

`not_my_key` is ensured both when the cell is empty (green cell in this example), or when the cell is busy, but occupied by an other key (a busy cell is in red in this example).



3.1.2 Impact of the modifications

The modifications have two main impacts: first, the accesses are not performed in the same order. The second impact is that the specification is not true anymore in the general case.

Access order: As explain above, with double hashing, cells of the map are not accessed by loop anymore. Hence, the `by_loop_for_all` lemma doesn't apply anymore. Let *stripe* be the function which, given a loop iteration, returns the index of the cell looked-up at this iteration (parametrized by *start*, *step* and *capacity*). This problem is solved by computing the antecedant of each cell.


Hence, at iteration *i*, for any cell `map[index]`, if the antecedent of `map[index]` is less than *i*, then `not_my_key(index)` is ensured.

The new way to ensure `not_my_key` for all indexes is then to ensures that all index has an antecedant w.r.t. the *stripe* function.

New requirements: However, it is not always the case that every cell is reached. An example is provided in Example 3. Actually, after the Chinese remainder theorem, the *step* and the *capacity* must be coprime in order to ensure that every cell is eventually tested.

Hence, this coprimeness is a new requirement of the specification. Technically, it is sufficient to have a capacity being a power of 2 and to have only odd *offset* hashes.

Example 3



3.2 The stripe_l_fp fixpoint

3.2.1 Definition

First, a fixpoint is defined, which returns the index to be updated after n iterations with an offset of $step$, starting from $start$ with capacity $capa$.

Definition 1: stripe(int start, int step, nat n, int capa)

```
fixpoint int stripe(int start, int step, nat n, int capa) {
  switch(n) {
    case zero: return start;
    case succ(m): return
      (stripe(start, step, m, capa) + step) % capa;
  }
}
```

The **stripe_l_fp** fixpoint builds a `list<option<nat>>` given a starting point, an offset, a number of accesses and a capacity. The base case of this fixpoint is to generate a `list` containing only `nones` (fixpoint `gen_none`), if `zero` accesses are performed. The recursive case is to update the $start + n * offset \% capa$ cell, using the above **stripe** fixpoint.

Definition 2: stripe_l_fp(int start, int step, nat bound, int capa)

A lemma ensuring that $stripe = start + n * step \% capa$ is proved.

The update (index, elem, list) fixpoint returns list with the index-th element updated to elem.

```

fixpoint list<option<nat>> stripe_l_fp(int start ,
  int step, nat n, int capa)
{
  switch(n) {
    case zero: return gen_none(nat_of_int(capa));
    case succ(m): return
      update(stripe(start , step , n, capa), some(n),
        stripe_l_fp(start , step , m, capa));
  }
}

```

Example 4: `stripe_l_fp(0, 2, 5, 7)`

Calling `stripe_l_fp(0, 2, 5, 7)` produces the following list. Notice that the base case returns a list containing only **none**s, not a list containing a **some(0)**.

none
some(4)
some(1)
some(5)
some(2)
none
some(3)

3.2.2 Properties

The main property required is that the number of cell containing **some(i)** (for any *i*) is equal to the number of steps done. This property will later be used to ensures that all cells are eventually reached (see Subsection 3.2.4). The function that count the number of such cells is named `count_some(list<option<nat>> list)`.

There are also other properties which are used internally to prove the `count.some` property. The main lemma is named `stripe_l`.

Lemma 1: Prototype of `stripe_l`

```
lemma list <option<nat> > stripe_l(int start, int step, nat n,
    int capa)
requires 0 <= start && start < capa && step > 0
    && n <= capa && coprime(step, capa) && step < capa;
ensures count_some(result) == n
    && length(result) == capa
    && true == up_to(nat_of_int(capa),
        (list_contains_stripes)(result, start, step))
    && true == up_to(nat_of_int(capa),
        (lst_opt_less_than_n)(result, n))
    && true == forall(result, opt_not_zero)
    && result == stripe_l_fp(start, step, n, capa)
    && coprime(step, capa);
```

`list_contains_stripe` ensures that if a cell contains `some(i)`, then `i` is the antecedant of the cell index w.r.t. `stripe`.

3.2.3 Proof of `stripe_l`

The proof of these properties relies on the fact that the same cell is not updated twice. Once this is ensured, the construction of the fixpoint ensures the validity of the properties.

Algorithm 1 shows the main steps of the proof. In the base case, all trivially holds. In the inductive case, if the `stripe(start, step, n, capa)`-th cell (i.e. the one hit at n -th iteration) already contains `some(i)`, the `list_contains_stripes` property ensures that `stripe(start, step, i, capa)`-th cell is the one we are hitting. Hence, $start + step \times i \% capa = start + step \times n \% capa$, with $n - i < capa$. Then the Chinese remainder theorem leads to a contradiction.

3.2.4 From `stripe` fixpoint to R

3.3 Proof of the *Chinese remainder theorem*

3.3.1 Properties

The goal of the *Chinese remainder theorem* is to highlight a contradiction in the `stripe_l` proof. We have that $diff \times step \% capa = 0$. The contradiction we want to highlight is that in the given environment, `diff` can only be 0, i.e. the supposed previous value is the same that the one we want to write, that is the n -th iteration is supposed to be already written to the list.

This is reduced to the following lemma: if $x \% n1 = 0$, $x \% n2 = 0$, $n1$ and $n2$ are coprime, and $x < n1 \times n2$, then $x = 0$. It is also required that $n1 > 0$, $n2 > 0$ and $x \geq 0$.

In `stripe_l`, this is applied to $x = diff \times step$, $n1 = step$ and $n2 = capa$.

This lemma is a direct consequence of the *uniqueness* property of the *Chinese remainder theorem*. Although it is simple to show informally, Verifast first requires to build `gcd` which is quite long. All the proof is done in a separate file `chinese_remainder_th.gh`. The proof

Input: int start, int step, nat n, int capa

```

switch n do
  case zero
    | // All hold by construction
  end
  case succ(m)
    // Recursive call, the termination is ensured by n > m
    list lst ← stripe_l(start, step, m, capa)
    // Now, we want to update the stripe(start, step, n, capa)-th to
    some(n)
    // Proof by contradiction that the cell contains none
    switch nth(stripe(start, step, n, capa), lst) do
      case some(i)
        assert start + i × step % capa = start + n × step % capa;
        assert (n - i) × step % capa = 0;
        assert n - i < capa;
        // The chinese remainder theorem applies and shows a
        contradiction
        chinese_remainder_theorem(step, capa, (n - i) × step);
      end
      case none
      end
    endsw
    // We now that the stripe(start, step, n, capa)-th cell contains
    a none, which we update to some(n), so the properties hold
    for the updated list.
    return update(stripe(start, step, n, capa), some(n), lst)
  end
endsw

```

$n - i$ is noted *diff*
in the following
parts

Algorithm 1: Proof of stripe_l

takes around 1000 lines of code (which less than 200 are *assert*-s or comments and could be remove).

3.3.2 Computation of *gcd*

3.3.3 *gcd* properties

3.3.4 Proof by contradiction

In Verifast, the proof of the `bin_chinese_remainder_theorem` lemma is quite long (approx. 300 lines). However, most of it is only arithmetic statements. Hence, informally, the proof is much shorter. Algorithm 2 sketches the main cases. The main part (*if* $x > 1$ branch) decompose x into $n1 \times k1 = n2 \times k2$. After justifying why $k1 \% n2 \neq 0$, it considers $\text{gcd}(k1, n2) = a$, which can not be 1. Then remaining case ($a \neq 1, k1 \% n2 \neq 0, x > 1$) calls recursively the theorem, on $\frac{n2}{a} = b$.

3.3.5 Assumed lemma

One lemma remains assumed:

Lemma 2: `gcd_mul`

```
lemma void gcd_mul(int n1, int n2, int n3)
  requires coprime(n1, n3) &*& coprime(n2, n3);
  ensures coprime(n1*n2, n3) &*& coprime(n1, n3)
    &*& coprime(n2, n3);
```

4 Conclusion

4.1 Validity of benchmark

4.2 Forthcoming work


```

if  $x = 1$  then
   $x \% n1 = 0 \Rightarrow n1 = 1$ ;
   $x \% n2 = 0 \Rightarrow n2 = 1$ ;
  assert  $n1 \times n2 = 1$ ;
  assert  $x = n1 \times n2$ ;
  contradiction;
else if  $x > 1$  then
   $x \% n1 = 0 \Rightarrow \exists k1 | n1 \times k1 = x$  ;
   $x \% n2 = 0 \Rightarrow \exists k2 | n2 \times k2 = x$  ;
  assert  $k1 \neq 0$ ;
  if  $k1 \% n2 = 0$  then
     $\beta \leftarrow k1 / n2$ ;
    assert  $\beta \times n2 = k1$ ;
    assert  $\beta \geq 1$ ;
    assert  $\beta \times n2 \leq k1$ ;
    assert  $x = n1 \times k1 \geq n1 \times n2$ ;
    contradiction;
  else
     $a \leftarrow gcd(k1, n2)$ ;
     $b \leftarrow n2 / a$ , assert  $b \neq 0$ ;
     $\gamma \leftarrow k2 / a$ ;
    assert  $gcd(b, \gamma) = 1$ ;
    if  $gcd(n1, b) \neq 1$  then
      assert  $gcd(n1, a \times b) \neq 1$ ;
      assert  $gcd(n1, n2) \neq 1$ ;
    end
    if  $a = 1$  then
      assert  $\gamma = k1 \wedge b = n2$ ;
       $gcd(b, \gamma) = 1 \wedge gcd(n1, b) = 1 \Rightarrow gcd(n1 \times \gamma, b) = 1$ ;
      contradiction  $gcd(x, n2) = 1$ ;
    else
      // The termination is ensured by  $b < n2$ 
       $bin\_chinese\_remainder\_theorem(n1, b, k2 \times b)$ ;
      assert  $k2 \times b = 0$ ;
      assert  $k2 = 0$ ;
      contradiction  $n2 \times k2 = x = 0$ ;
    end
  end
end
else
  assert  $x = 0$ ;
end

```

Algorithm 2: Proof of `bin_chinese_remainder_theorem`

```

int i = 0;
for (; i < capacity; ++i)
  /*@ invariant ...  $\mathcal{E}\mathcal{E}$ 
    true == up_to(nat_of_int(i),
      (byLoopNthProp)(ks, (not_my_key)(k),
        capacity, start));
  @*/
  /*@ decreases capacity - i;
  {
    int index = loop(start + i, capacity);
    int bb = busybits[index];
    int kh = k_hashes[index];
    void* kp = keyps[index];
    if (bb != 0 && kh == key_hash) {
      if (eq(kp, keyp)) {
        /*@ hmap_find_this_key(hm, index, k);
        return index;
      }
    } else {
      /*@ if (bb != 0) no_hash_no_key(ks, khs, k, index, hsh);
      /*@ if (bb == 0) no_bb_no_key(ks, bbs, index);
    }
    /*@ assert(true == not_my_key(k, nth(index, ks)));
  }
  /*@ by_loop_for_all(ks, (not_my_key)(k),
    start, capacity, nat_of_int(capacity));
  @*/
  /*@ assert true == up_to(nat_of_int(length(ks)),
    (nthProp)(ks, (not_my_key)(k)));
  @*/
  /*@ no_key_found(ks, k);

```

Figure 1: Original *for*-loop for searching a key